

Main.py

```
1  # coding: utf-8
2  import argparse
3  import time
4  import math
5  import os
6  import torch
7  import torch.nn as nn
8
9  import data
10 import model
11
12 parser = argparse.ArgumentParser(description='PyTorch Wikitext-2 RNN/LSTM/GRU/Transformer Lan
13 parser.add_argument('--data', type=str, default='./data/wikitext-2',
14                     help='location of the data corpus')
15 parser.add_argument('--model', type=str, default='LSTM',
16                     help='type of network (RNN_TANH, RNN_RELU, LSTM, GRU, Transformer)')
17 parser.add_argument('--emsize', type=int, default=200,
18                     help='size of word embeddings')
19 parser.add_argument('--nhid', type=int, default=200,
20                     help='number of hidden units per layer')
21 parser.add_argument('--nlayers', type=int, default=2,
22                     help='number of layers')
23 parser.add_argument('--lr', type=float, default=20,
24                     help='initial learning rate')
25 parser.add_argument('--clip', type=float, default=0.25,
26                     help='gradient clipping')
27 parser.add_argument('--epochs', type=int, default=40,
28                     help='upper epoch limit')
29 parser.add_argument('--batch_size', type=int, default=20, metavar='N',
30                     help='batch size')
31
32 parser.add_argument('--bptt', type=int, default=35,
33                     help='sequence length')
34 parser.add_argument('--dropout', type=float, default=0.2,
35                     help='dropout applied to layers (0 = no dropout)')
36 parser.add_argument('--tied', action='store_true',
37                     help='tie the word embedding and softmax weights')
38 parser.add_argument('--seed', type=int, default=1111,
39                     help='random seed')
40 parser.add_argument('--cuda', action='store_true',
41                     help='use CUDA')
42 parser.add_argument('--log-interval', type=int, default=200, metavar='N',
43                     help='report interval')
44 parser.add_argument('--save', type=str, default='model.pt',
45                     help='path to save the final model')
46 parser.add_argument('--nhead', type=int, default=2,
47                     help='the number of heads in the encoder/decoder of the transformer model')
48 parser.add_argument('--dry-run', action='store_true',
49                     help='verify the code and the model')
50
51 args = parser.parse_args()
52
53 # Set the random seed manually for reproducibility.
54 torch.manual_seed(args.seed)
55 if torch.cuda.is_available():
56     if not args.cuda:
57         print("WARNING: You have a CUDA device, so you should probably run with --cuda.")
58
59 device = torch.device("cuda" if args.cuda else "cpu")
```

```

59 #####
60 # Load data
61 #####
62
63 corpus = data.Corpus(args.data)
64 # Starting from sequential data, batchify arranges the dataset into columns.
65 # For instance, with the alphabet as the sequence and batch size 4, we'd get
66 # [ a g m s ]
67 # [ b h n t ]
68 # [ c i o u ]
69 # [ d j p v ]
70 # [ e k q w ]
71 # [ f l r x ].
72 # These columns are treated as independent by the model, which means that the
73 # dependence of e. g. 'g' on 'f' can not be learned, but allows more efficient
74 # batch processing.
75
76 def batchify(data, bsz):
77     # Work out how cleanly we can divide the dataset into bsz parts.
78     nbatch = data.size(0) // bsz
79     # Trim off any extra elements that wouldn't cleanly fit (remainders).
80     data = data.narrow(0, 0, nbatch * bsz)
81     # Evenly divide the data across the bsz batches.
82     data = data.view(bsz, -1).t().contiguous()
83     return data.to(device)
84
85 eval_batch_size = 10
86 train_data = batchify(corpus.train, args.batch_size)
87 val_data = batchify(corpus.valid, eval_batch_size)
88 test_data = batchify(corpus.test, eval_batch_size)

```

```

90 #####
91 # Build the model
92 #####
93
94 ntokens = len(corpus.dictionary)
95 if args.model == 'LSTM':
96     model = model.RNNModel(args.model, ntokens, args.emsize, args.nhid, args.nlayers, args.dr
97
98 criterion = nn.CrossEntropyLoss()
99
100 #####
101 # Training code
102 #####
103
104 def repack_hidden(h):
105     """Wraps hidden states in new Tensors, to detach them from their history."""
106
107     if isinstance(h, torch.Tensor):
108         return h.detach()
109     else:
110         return tuple(repackage_hidden(v) for v in h)
111
112
113 # get_batch subdivides the source data into chunks of length args.bptt.
114 # If source is equal to the example output of the batchify function, with
115 # a bptt-limit of 2, we'd get the following two Variables for i = 0:
116 # [ a g m s ] [ b h n t ]
117 # [ b h n t ] [ c i o u ]
118 # Note that despite the name of the function, the subdivision of data is not
119 # done along the batch dimension (i.e. dimension 1), since that was handled

```

```
120 # by the batchify function. The chunks are along dimension 0, corresponding
121 # to the seq_len dimension in the LSTM.
```

```
122
123 def get_batch(source, i):
124     seq_len = min(args.bptt, len(source) - 1 - i)
125     data = source[i:i+seq_len]
126     target = source[i+1:i+1+seq_len].view(-1)
127     return data, target
128
129
130 def evaluate(data_source):
131     # Turn on evaluation mode which disables dropout.
132     model.eval()
133     total_loss = 0.
134     ntokens = len(corpus.dictionary)
135     hidden = model.init_hidden(eval_batch_size)
136     with torch.no_grad():
137         for i in range(0, data_source.size(0) - 1, args.bptt):
138             data, targets = get_batch(data_source, i)
139             output, hidden = model(data, hidden)
140             hidden = repackage_hidden(hidden)
141             total_loss += len(data) * criterion(output, targets).item()
142     return total_loss / (len(data_source) - 1)
```

```
143
144
145 def train():
146     # Turn on training mode which enables dropout.
147     model.train()
148     total_loss = 0.
149     start_time = time.time()
150     ntokens = len(corpus.dictionary)
151     hidden = model.init_hidden(args.batch_size)
152     for batch, i in enumerate(range(0, train_data.size(0) - 1, args.bptt)):
153         data, targets = get_batch(train_data, i)
154         # Starting each batch, we detach the hidden state from how it was previously produced
155         # If we didn't, the model would try backpropagating all the way to start of the datas
156         model.zero_grad()
157         hidden = repackage_hidden(hidden)
158         output, hidden = model(data, hidden) # model.forward(data, hidden)
159         loss = criterion(output, targets)
160         loss.backward()
161         # `clip_grad_norm` helps prevent the exploding gradient problem in RNNs / LSTMs.
162         torch.nn.utils.clip_grad_norm_(model.parameters(), args.clip)
163
164         for p in model.parameters():
165             p.data.add_(p.grad, alpha=-lr)
166             # optimizer: admw adafactor ...
167     total_loss += loss.item()
```

```
168
169     if batch % args.log_interval == 0 and batch > 0:
170         cur_loss = total_loss / args.log_interval
171         elapsed = time.time() - start_time
172         print('| epoch {:3d} | {:5d}/{:5d} batches | lr {:02.2f} | ms/batch {:5.2f} | '
173               'loss {:5.2f} | ppl {:8.2f}'.format(
174             epoch, batch, len(train_data) // args.bptt, lr,
175             elapsed * 1000 / args.log_interval, cur_loss, math.exp(cur_loss)))
176         total_loss = 0
177         start_time = time.time()
178     if args.dry_run:
179         break
180
181 # Loop over epochs.
182 lr = args.lr
183 best_val_loss = None
```

```

184
185 # At any point you can hit Ctrl + C to break out of training early.
186
187 try:
188     for epoch in range(1, args.epochs+1):
189         epoch_start_time = time.time()
190         train()
191         val_loss = evaluate(val_data)
192         print('-' * 89)
193         print('| end of epoch {:3d} | time: {:5.2f}s | valid loss {:5.2f} | '
194               | 'valid ppl {:8.2f}'.format(epoch, (time.time() - epoch_start_time),
195                                           val_loss, math.exp(val_loss)))
196         print('-' * 89)
197         # Save the model if the validation loss is the best we've seen so far.
198         if not best_val_loss or val_loss < best_val_loss:
199             with open(args.save, 'wb') as f:
200                 torch.save(model, f)
201                 best_val_loss = val_loss
202         else:
203             # Anneal the learning rate if no improvement has been seen in the validation data
204             lr /= 4.0
205 except KeyboardInterrupt:
206     print('-' * 89)
207     print('Exiting from training early')

```

```

210 # Load the best saved model.
211 with open(args.save, 'rb') as f:
212     model = torch.load(f)
213     # after load the rnn params are not a continuous chunk of memory
214     # this makes them a continuous chunk, and will speed up forward pass
215     # Currently, only rnn model supports flatten_parameters function.
216     if args.model in ['RNN_TANH', 'RNN_RELU', 'LSTM', 'GRU']:
217         model.rnn.flatten_parameters()
218
219 # Run on test data.
220 test_loss = evaluate(test_data)
221 print('=' * 89)
222 print('| End of training | test loss {:5.2f} | test ppl {:8.2f}'.format(
223       test_loss, math.exp(test_loss)))
224 print('=' * 89)

```

Data.py

```

1  import os
2  from io import open
3  import torch
4
5  class Dictionary(object):
6      def __init__(self):
7          self.word2idx = {}
8          self.idx2word = []
9
10     def add_word(self, word):
11         if word not in self.word2idx:
12             self.idx2word.append(word)
13             self.word2idx[word] = len(self.idx2word) - 1
14             return self.word2idx[word]
15
16     def __len__(self):
17         return len(self.idx2word)

```

```

20 class Corpus(object):
21     def __init__(self, path):
22         self.dictionary = Dictionary()
23         self.train = self.tokenize(os.path.join(path, 'train.txt'))
24         self.valid = self.tokenize(os.path.join(path, 'valid.txt'))
25         self.test = self.tokenize(os.path.join(path, 'test.txt'))
26
27     def tokenize(self, path):
28         """Tokenizes a text file."""
29         assert os.path.exists(path)
30         # Add words to the dictionary
31         with open(path, 'r', encoding="utf8") as f:
32             for line in f:
33                 words = line.split() + ['<eos>']
34                 for word in words:
35                     self.dictionary.add_word(word)
36
37         # Tokenize file content
38         with open(path, 'r', encoding="utf8") as f:
39             idss = []
40             for line in f:
41                 words = line.split() + ['<eos>']
42                 ids = []
43                 for word in words:
44                     ids.append(self.dictionary.word2idx[word])
45                 idss.append(torch.tensor(ids).type(torch.int64))
46             ids = torch.cat(idss)
47
48         return ids

```

Model.py

```

1  import math
2  import torch
3  import torch.nn as nn
4  import torch.nn.functional as F
5
6  class RNNModel(nn.Module):
7      """Container module with an encoder, a recurrent module, and a decoder."""
8
9      def __init__(self, rnn_type, ntoken, ninp, nhid, nlayers, dropout=0.5, tie_weights=False):
10         super(RNNModel, self).__init__()
11         self.ntoken = ntoken
12         self.encoder = nn.Embedding(ntoken, ninp)
13         if rnn_type in ['LSTM', 'GRU']:
14             self.rnn = getattr(nn, rnn_type)(ninp, nhid, nlayers, dropout=dropout)
15         else:
16             try:
17                 nonlinearity = {'RNN_TANH': 'tanh', 'RNN_RELU': 'relu'}[rnn_type]
18             except KeyError:
19                 raise ValueError( """An invalid option for `--model` was supplied,
20                                options are ['LSTM', 'GRU', 'RNN_TANH' or 'RNN_RELU']""")
21             self.rnn = nn.RNN(ninp, nhid, nlayers, nonlinearity=nonlinearity, dropout=dropout)
22         self.decoder = nn.Linear(nhid, ntoken)
23
24         self.drop = nn.Dropout(dropout)
25
26         self.init_weights()
27
28         self.rnn_type = rnn_type
29         self.nhid = nhid
30         self.nlayers = nlayers

```

